

MARKOVIAN

markovian.co.uk

Company Overview

AI Compliance Consulting for the Age of Regulated AI

UK/EU AI Compliance Advisory

GDPR · EU AI Act · SOC 2 · ISO/IEC 42001

May 2026

Mat 20/05/26

1. Executive Summary

What We Do

Markovian is a Cambridge-based AI compliance consultancy that helps small and medium-sized technology companies navigate the regulations that govern how artificial intelligence can be built and deployed. We exist because the EU AI Act, GDPR, and emerging UK frameworks are now legally binding — and the companies that need the most help are the ones with the least guidance.

The regulatory landscape for AI has shifted fundamentally. The EU AI Act entered into force in August 2024, prohibited AI practices applied from February 2025, and full enforcement of high-risk AI obligations begins in August 2026. Organisations deploying AI that touches hiring, lending, healthcare, education, or public services face serious legal exposure if they cannot demonstrate compliance. Fines reach up to 7% of global annual turnover.

The market is underserved at the SME level. Large consultancies serve only enterprise clients at prohibitive cost. Compliance automation tools like Sprinto handle generic IT controls but have no capacity to evaluate AI-specific risks — they cannot assess whether a model violates Article 22 of GDPR, classify a system under the EU AI Act, or write the technical documentation regulators require. No specialist, technically credible, SME-focused AI compliance firm exists in the UK today.

Markovian fills that gap. Our team combines genuine AI and engineering expertise (Cambridge Electrical Engineering, UCL Statistics, Cambridge Information Engineering) with direct industry experience in financial services AI. We offer two complementary services: structured compliance audits for companies with existing AI systems, and an end-to-end compliance partnership for companies building new AI projects from scratch. Both services are anchored by our proprietary Markovian AI Risk Framework (MARF), a structured scoring model that identifies where a client's greatest regulatory and safety exposure lies and prioritises what to address first.

We are not a slideshow factory. We are a lean, technically sophisticated advisory firm with a repeatable methodology — one that is designed to scale without growing into a large headcount consultancy.

2. The Problem

2.1 The Regulatory Moment Has Arrived

For years, AI regulation was theoretical. It is no longer. The EU AI Act is now in force with a phased enforcement schedule that peaks in August 2026, when the majority of obligations — including all high-risk AI system requirements — become enforceable. GDPR, which has applied since 2018, contains provisions (particularly Article 22 on automated decision-making) that are now being actively enforced against AI systems. ISO/IEC 42001, the first international AI management system standard, was published in December 2023 and is being adopted as a governance benchmark by enterprise procurement teams and government bodies.

In the UK, while no dedicated AI Act exists yet, the Data (Use and Access) Act 2025 introduced the first statutory AI-relevant obligations, and a comprehensive AI Bill is expected in 2026. UK

companies serving EU customers are simultaneously in scope for the EU AI Act regardless of where they are incorporated.

2.2 The Gap at the SME Level

Large enterprises can hire internal legal and compliance teams or engage top-tier consultancies. SMEs and startups — the fastest-growing adopters of AI — cannot. The gap they face is not just about cost. It is about comprehension. The companies most urgently in need of compliance help are the ones with no internal expertise to even evaluate what they do not know.

This insight came directly from experience. Ben, Markovian’s co-founder and lead AI compliance specialist, worked within the AI safety and development team at GB Bank, where he repeatedly observed smaller FinTech partners and clients attempting to build AI-driven products with no understanding of the regulatory framework governing them. The question was never “how do we comply” — it was “we didn’t know any of this applied to us.”

2.3 What Existing Solutions Cannot Do

Current market offerings address different problems:

Type	What they do	What they cannot do
Sprinto / Vanta / Drata	Automate SOC 2 / ISO 27001 evidence collection via cloud integrations	Assess AI-specific risk, evaluate models, classify EU AI Act tiers, or write technical documentation
Large consultancies (Accenture, Capco)	Enterprise AI strategy and transformation at scale	Serve SMEs; prohibitively expensive; often lack deep technical AI model expertise
Law firms	Provide legal opinions on regulatory exposure	Understand how models work, evaluate training data, or give technically grounded implementation guidance
Exona Lab	Help insurers price and underwrite AI risk	Help companies comply — they serve insurers, not the companies deploying AI

3. Why Now

Three forces have converged to create the window Markovian is built for:

Regulatory deadlines are imminent and real

- February 2025: EU AI Act prohibitions on unacceptable-risk AI took effect
- August 2025: GPAI model obligations and EU AI Act governance rules applied
- August 2026: Full high-risk AI system compliance obligations enforce, with fines up to 7% of global turnover
- The UK Data (Use and Access) Act 2025 introduced the first statutory AI-relevant UK obligations

- A comprehensive UK AI Bill is expected in 2026, extending obligations further

SME AI adoption is accelerating ahead of compliance readiness

- The fastest-growing AI adopters are FinTechs, LegalTechs, HR platforms, and healthtech startups
- These companies are deploying AI into hiring, lending, diagnostics, and legal decisions — all high-risk categories under the EU AI Act — without compliance infrastructure
- Enterprise customers and investors are now requiring compliance proof before signing contracts

ISO/IEC 42001 creates a new certification market

- Published December 2023, ISO/IEC 42001 is the first international AI management system standard
- UK government guidance explicitly recommends it as a governance benchmark
- Enterprise procurement and government contracts are beginning to require it
- No widespread certification readiness service exists for SMEs yet

4. Our Solution

Markovian delivers AI compliance through two service lines and one proprietary methodology.

The Markovian AI Risk Framework (MARF)

The Markovian AI Risk Framework (MARF) is our proprietary risk scoring model, used internally at the start of every engagement. It takes structured inputs about a client’s AI system — what it decides, what data it processes, who it affects, how it is deployed — and produces a weighted risk map across five dimensions. The output identifies the top priority areas of regulatory and safety exposure, sequences remediation by impact and effort, and directly informs the technical documentation required by both the EU AI Act and ISO/IEC 42001.

4.1 MARF — The Five Risk Dimensions

Every engagement begins with MARF. The five axes are:

Dimension	What it measures	Regulatory anchor
Regulatory Exposure	Which specific laws apply and how severely, given sector and geography	EU AI Act risk tier, GDPR, UK DUA Act
Decision Impact	Severity and reversibility of harm if the AI produces a wrong output	EU AI Act Art. 6 (high-risk), GDPR Art. 22
Data Sensitivity	Classification and volume of personal data processed, including special categories	GDPR, ISO/IEC 42001 Section A.7
Technical Opacity	How explainable and auditable the model is; ability to satisfy transparency obligations	GDPR recital 71, EU AI Act transparency, ISO 42001 B.6

Governance Maturity	Existing documentation, oversight, and process; gap between current state and required state	ISO/IEC 42001, SOC 2 TSC, EU AI Act Art. 9
---------------------	--	--

The output is a risk tier (Critical / High / Medium / Low), a prioritised list of the top three to five vulnerabilities, and a recommended action plan ordered by regulatory urgency and implementation effort. Clients who complete this process are simultaneously building their ISO/IEC 42001 risk documentation — a material side benefit that reduces time to certification.

4.2 Service Line 1: AI Compliance Audit

For companies with an existing AI system or one approaching deployment.

- The client brings an existing or near-complete AI system
- We run MARF as a diagnostic — an MRI scan of what has already been built
- We review documentation, conduct technical interviews with developers, and where relevant examine model outputs and data pipelines
- We produce a gap analysis report: current state vs. required state under applicable regulations
- We deliver a prioritised remediation roadmap with concrete, technically-grounded recommendations
- Engagement length: typically two to four weeks
- Delivery: primarily remote, with an initial kick-off and final presentation meeting

Typical trigger events: EU AI Act compliance deadlines, Series A investor due diligence, enterprise customer security questionnaires, SOC 2 audit preparation, or responding to a regulator query.

4.3 Service Line 2: Compliant AI by Design

For companies that want to build an AI project but need to do it safely and compliantly from the first line of architecture.

This is the service Ben identified as the critical gap at GB Bank. Smaller FinTech partners were not failing to comply with regulations they knew about — they were building systems with no awareness that the regulations existed at all. The consequence is expensive: retrofitting compliance into a deployed system is dramatically harder than designing it in from the start.

- The client comes with an idea or early-stage project, not a finished system
- We begin with a discovery workshop to understand the business problem, the proposed AI solution, the data available, and the intended user population
- We run MARF as a design guide — the risk scores become design requirements for the development team
- We remain engaged throughout the project lifecycle: requirements, data strategy, model selection, documentation, testing, deployment, and monitoring
- We translate regulatory requirements into concrete engineering steps the client's developers can act on directly
- Engagement length: three to twelve months depending on project complexity
- Delivery: blended — in-person workshops at key milestones, regular remote check-ins, ongoing advisory access

This service line creates the deepest client relationships and the highest switching costs. Once Markovian is embedded in a client's development process, we hold institutional knowledge of their systems that cannot easily be replicated by a new entrant. It is also the service that most directly advances AI safety: companies that build with compliance in mind from the start deploy systems that are more transparent, more auditable, and less likely to cause harm.

Our Value Proposition

We are not selling compliance paperwork. We are selling the ability to build AI confidently — knowing that what you are building will not be a liability the moment it is deployed.

5. Target Market

5.1 Primary Sectors

We focus initial efforts on three sectors where regulatory exposure is highest, AI adoption is accelerating, and SME compliance infrastructure is most absent:

FinTech and Challenger Banking — Primary Beachhead

- AI is embedded in credit scoring, fraud detection, KYC, and automated customer decisions — all high-risk under the EU AI Act
- Smaller FinTechs frequently lack any internal compliance expertise for AI
- Ben's direct experience at GB Bank gives Markovian immediate credibility and a warm network in this sector
- FCA-regulated firms face compounding obligations: GDPR, EU AI Act, FCA model risk guidance, and Consumer Duty

LegalTech

- AI tools used in contract review, case prediction, due diligence, and sentencing analysis are explicitly classified as high-risk
- LegalTech startups face professional liability if their AI tools produce non-compliant outputs
- High willingness to pay: a single compliance failure can expose the vendor and their law firm customers simultaneously

HR and Recruitment Technology

- The EU AI Act explicitly classifies recruitment, candidate selection, performance evaluation, task allocation, and worker monitoring as high-risk AI use cases
- This is one of the largest categories of enterprise AI deployment and one of the least compliant
- Every company selling an AI hiring tool into the EU market faces this obligation from August 2026

5.2 Secondary Sectors

- HealthTech and MedTech — high data sensitivity, MHRA reviewing AI regulation, AI diagnostics classified as high-risk
- EdTech — AI in education explicitly high-risk under EU AI Act; growing sector with limited compliance resources
- Government contractors and public sector vendors — high value engagements; companies bidding for AI-related government contracts must demonstrate compliance

5.3 Client Profile

The ideal Markovian client is a company with ten to two hundred employees that:

- Is building or has built an AI system used in a high-stakes decision context
- Operates in the UK or EU, or serves customers in those jurisdictions
- Has no internal AI governance or compliance function
- Is approaching a regulatory deadline, enterprise procurement requirement, or fundraising event that requires compliance evidence
- Has developers and a CTO who speak the language of engineering, not legal

6. Competitive Position and Differentiation

Markovian's differentiation rests on four properties that are structurally difficult for existing players to replicate:

1. Technical depth that legal and audit backgrounds cannot match

A Cambridge Electrical Engineer, a UCL Statistics graduate with direct FinTech AI experience, and a Cambridge Information Engineer with AI safety research background form a team that can read a model card, evaluate a training pipeline, assess algorithmic fairness, and explain how a model's outputs could violate GDPR — in the same conversation as a developer. No traditional compliance consultancy can do this. No SaaS tool can do this. This is the core of our value.

2. AI-specific scope that generic compliance tools do not cover

Sprinto, Vanta, and similar tools automate the collection of evidence for SOC 2 and ISO 27001 IT controls. They cannot classify an AI system's risk tier, evaluate whether a model's decision logic satisfies transparency obligations, assess training data for bias, or write the technical documentation the EU AI Act mandates for high-risk systems. These are not gaps in their roadmap — they are outside their business model. Markovian operates in the layer above them, and can complement their infrastructure rather than compete with it.

3. Proximity to the problem through direct industry experience

The insight that drives Markovian came from inside the problem, not from a market analysis. Ben's direct observation of smaller FinTechs attempting to build AI systems with no compliance awareness gives us an authentic understanding of the client's perspective and a network of warm contacts. This is not a market we identified from the outside — it is a problem we watched unfold from the inside.

4. A proprietary methodology that makes engagements repeatable and scalable

The MARF model means every engagement starts from a structured, defensible analytical foundation rather than a blank page. This reduces the time per engagement, increases consistency of output, and creates an intellectual property asset that grows in value as we validate it across more clients. It is also the foundation for a future SaaS product if and when the team chooses to productise.

7. The Team

Ben Co-Founder

UCL Statistics graduate. Former prediction modelling specialist at Statz.ai. AI safety and development team member at GB Bank, where he worked under senior leadership on AI compliance for a newly certified challenger bank. Direct first-hand insight into the compliance gap facing smaller FinTech and banking clients. The technical and domain anchor of Markovian's services.

Mat Co-Founder

Cambridge Electrical Engineering graduate. Currently on GPU engineering internship at Qualcomm. Responsible for company strategy, client relationships, business development, and operations. Background in engineering systems gives practical technical grounding for translating compliance requirements into language developers understand.

Sassan AI Safety & Research

Cambridge Information Engineering student with AI specialisation modules. Deep interest and background in AI safety research. Strong command of the AI safety ecosystem, funding landscape, and research frontier. Responsible for technical AI safety framing, research partnerships, and ecosystem positioning. Introduced the team to BlueDot Impact and the broader effective AI safety community.

8. Business Model

8.1 Revenue Model

Markovian operates an advisory model with fixed-scope, clearly priced engagements. We do not sell time by the hour — we sell outcomes. This keeps overhead low, makes costs predictable for clients, and forces us to design repeatable processes rather than reinventing each engagement from scratch.

Product	Price range	Engagement type
AI Risk Snapshot (MARF diagnostic report)	£1,500 – £2,500	Fixed price, 3–5 hours of analyst time, 1–2 week turnaround
AI Compliance Audit	£5,000 – £15,000	Fixed scope, 2–4 weeks, gap analysis + remediation roadmap
Compliant AI by Design	£3,000 – £7,000 / month	Retainer, 3–12 months, embedded advisory through project lifecycle

ISO/IEC 42001 Readiness	£8,000 – £18,000	Fixed scope, structured path to certification readiness
-------------------------	------------------	---

The AI Risk Snapshot is the entry product — low commitment for the client, low overhead for us, and a natural conversion point into a full audit or design engagement when the client sees their risk score.

8.2 What We Will Not Do

Markovian will not write production code, build AI systems, or manage client IT infrastructure. Implementation is always carried out by the client's own development team or a referred implementation partner. This keeps us lean, objective, and out of scope creep. It also avoids liability for building or maintaining systems.

8.3 Referral Partnerships

For clients that need implementation support beyond our advisory scope, we maintain a curated list of trusted partners — specialist data engineers, privacy tech vendors, and security firms. Referral arrangements generate supplementary income while ensuring clients get end-to-end support.

9. Roadmap

Phase 0 — Foundation (Now – June 2026)

Immediate priorities before the first client engagement:

- Submit BlueDot Impact Incubator Week application by 26 May 2026
- Update markovian.co.uk to reflect the two service lines, the team, and the MARF concept
- Build MARF v1: a structured internal scoring model and accompanying client questionnaire
- Draft the AI Risk Snapshot template: the report format clients receive after a MARF assessment
- Activate Ben's FinTech and banking network: initial outreach to 20–30 warm contacts
- Identify five to ten target companies with obvious EU AI Act exposure and approaching deadlines

Phase 1 — First Clients (June – September 2026)

The EU AI Act's August 2026 enforcement date creates a concentrated window of urgency. Phase 1 is designed to capture it.

- Deliver two to four AI Risk Snapshots as paid diagnostic engagements
- Convert at least one into a full AI Compliance Audit
- Begin at least one Compliant AI by Design retainer engagement
- Validate MARF against real client systems and refine the scoring model
- Build the first case study: a named outcome ("Helped [company] achieve GDPR and EU AI Act readiness ahead of August 2026 enforcement")
- If BlueDot Incubator Week is successful: use funding to accelerate website, tooling, and first hire or associate

Phase 2 — Establish and Expand (October 2026 – March 2027)

- Compliant AI by Design becomes the flagship product and primary marketing message
- Expand into LegalTech and HR Tech sectors using FinTech case studies as proof
- Introduce ISO/IEC 42001 readiness as a standalone service, as enterprise demand for certification grows
- Consider first associate hire or contractor arrangement if demand exceeds core team capacity
- Explore partnerships with law firms and DPOs who encounter AI compliance questions beyond their expertise

Phase 3 — Productise (2027 onwards)

- MARF validated across 15–25 client engagements: explore productising as a licensed tool
- A lightweight SaaS version of MARF could be offered to DPOs, law firms, and smaller consultancies
- Recurring licence revenue reduces dependence on project-based income
- Ongoing regulatory monitoring service: regular updates to clients as the EU AI Act, UK AI Bill, and ISO standards evolve

10. AI Safety Framing

Markovian is not a compliance business that happens to work on AI. It is an AI safety business that operates through the mechanism of compliance.

The most dangerous AI systems are not the ones that fail in the lab. They are the ones deployed into consequential real-world decisions — hiring, lending, medical diagnosis, legal outcomes — without adequate testing, documentation, oversight, or transparency. The EU AI Act and related frameworks exist precisely because regulators recognised that market forces alone would not produce safe AI deployment.

The companies most likely to deploy unsafe AI are not the large labs with AI safety teams and legal departments. They are the hundreds of small and medium FinTechs, LegalTechs, and HR platforms that are embedding AI into decisions that affect people's lives, with no internal expertise to evaluate what they are doing. These companies are not malicious. They are uninformed. That is exactly the gap Markovian exists to close.

The AI Safety Case for Markovian

Making AI safety commercially accessible to the companies that need it most — not just the ones that can afford it — is how responsible AI deployment scales. Compliance is the market mechanism through which safety becomes a default, not an exception.

This framing also positions Markovian as complementary to, rather than competing with, other AI safety organisations. Exona Lab makes AI risks insurable. Policy organisations shape the regulatory framework. Research labs study alignment. Markovian helps the companies at the deployment end of the chain actually implement what all of those efforts produce. We are the last mile of responsible AI deployment.